

Amortized eFPGA Layout Search via Reinforcement Learning: Zero-Shot Generalization Across Circuit Benchmarks

Anonymous Author(s)

Abstract

Field-Programmable Gate Arrays (FPGAs) offer essential hardware flexibility through programmable logic and interconnects. Embedded FPGAs (eFPGAs) extend this capability by integrating reconfigurable fabrics directly into fixed-silicon Systems-on-Chip (SoCs) or ASICs, enabling post-fabrication functional updates and robust IP security. Hence, discovering optimized layouts for these reconfigurable fabrics is essential to fully realize their architectural potential. We train a single graph-and-grid-conditioned Reinforcement Learning (RL) policy across various benchmark circuits to place heterogeneous blocks (H-blocks, such as DSPs and BRAMs) and select the optimal fabric aspect ratio for the eFPGA architectures. We thereby generate layouts jointly optimized for Power, Performance, and Area (PPA). To demonstrate generalizability, we evaluate the trained RL policy on benchmark circuits completely unseen during training, producing highly efficient zero-shot layouts. Across three random seeds, the policy achieves up to a 58.7% zero-shot ADP reduction on unseen benchmark circuits relative to traditional island-style architectures.

Keywords

FPGA, Reinforcement Learning (RL), Graph Convolutional Networks (GCN), Convolutional Neural Networks (CNN), Macro Placement, Reconfigurable Computing

1 Introduction

A generic FPGA inherently over-provisions hardware resources for any single application, carrying extra routing and heterogeneous blocks (*H-blocks*: DSP slices, BRAMs) that the deployed netlist never fully utilizes. Sizing an eFPGA fabric for one specific netlist recovers much of this overhead. The product of routing area, critical-path delay, and total power (ADP throughout) directly quantifies this overhead; minimizing this metric benefits every eFPGA application [3, 19].¹

Prior efforts optimize fabric layouts on a per-design basis using computationally expensive search heuristics like Genetic Algorithms [19] or NSGA-II [3]. These approaches restart the search from scratch for every new circuit, transferring no knowledge across runs and incurring the full computational search cost each time.

We train a single graph-and-grid-conditioned RL policy to determine H-block placement and fabric aspect ratio, addressing a domain where computationally expensive per-instance GA search previously served as the only option. Unlike RL methods [4, 10] that assign blocks to a pre-fabricated grid, our policy directly constructs the fabric layout itself. After jointly training on 11 circuits evaluated via the standard VTR flow, we apply a single model checkpoint

zero-shot to unseen circuits, including fabrics up to $3.4\times$ the maximum training pool area. We train only once, after which evaluating a new circuit requires just a single policy rollout. We compare comprehensively against a from-scratch GA under an identical oracle, metric, and baseline and find that while the GA retains a marginal advantage in absolute ADP reduction for some circuits, the trained policy consistently achieves highly competitive results. Crucially, for larger circuits where evolutionary runtimes scale prohibitively, our zero-shot policy provides a superior alternative by delivering optimized layouts instantly without the expensive computational overhead (Section 5.2).

We contribute:

-
-
-
-

2 Background and Related Work

2.1 Per-Design eFPGA Layout

A generic island-style FPGA provisions H-blocks (DSP, BRAM) and routing for arbitrary workloads; any single application leaves much of this unused, paying area and delay for resources it never exercises. This motivates fitting tile placement and fabric dimensions to one specific netlist for a compact, efficient fabric. Applications include IP security (function absent from fixed silicon until configured [6, 20]), domain-specific acceleration, and post-fabrication reconfigurability [19]. Subsequent work reduces overhead by shrinking under-utilized routing [14] or de-regularizing the fabric [8], both motivating the non-island-style layouts we synthesize. All rely on per-design search; none amortizes layout effort across netlists. GOLDS [19] introduced layout search built for a single design: a GA over H-block placement with Area $\times$ Delay fitness through the real VTR flow. A follow-up [3] applies NSGA-II, also co-evolving CLB LUT size and H-block sizing. Both report per-design gains on their own two-term Area $\times$ Delay metric, a different quantity from the three-term ADP we report. Both start from a random population for every new design, transferring nothing from prior runs. We target the same H-block placement problem using an RL approach instead; incorporating block-sizing co-evolution [3] into the RL policy is a natural extension (Section 6).

2.2 Learned Policies for Physical Design

AlphaChip [17] trains an RL agent to place ASIC macros using a learned proxy and reports cross-design transfer: the amortization argument we make here, at a much larger scale. Subsequent work explores visual representations [16], joint placement-and-routing [5], and transformer policies [15], all supporting the claim that a trained policy amortizes layout cost across designs as per-instance search cannot. Unlike AlphaChip, we evaluate every candidate against the

¹GOLDS [19] uses the same abbreviation, ADP, for a different, two-term (logical_area $\times$ delay) product; the proposed framework employs a three-term (routing_area $\times$ delay $\times$ power) product throughout this paper.

true VTR flow, not a proxy: affordable at our scale, where a VTR run takes seconds to low minutes (Section 4.1).

Within the FPGA setting, RL has guided placement onto an *already-fabricated* grid (annealing perturbations in VPR [10] or a Gym-style VPR environment [4]), but always assigns blocks to a fixed grid, never shapes it. Transferable EDA policies exist for tool tuning [12] and global placement [11]; our policy instead synthesizes the fabric geometry.

2.3 VTR/VPR Background

We build on the open-source Verilog-to-Routing (VTR) [9]: Yosys synthesizes the design [7] and VPR [2] packs, places, and routes against an architecture XML of tile types, positions, and connectivity.

3 Methodology

3.1 Problem Formulation

A fabric optimized for one circuit trades general-purpose flexibility for a significantly smaller, more efficient implementation: it accommodates parameter changes, post-fabrication fixes, and a known set of operational modes on that one primary function, not arbitrary workloads with substantially different H-block requirements.

The *policy* configures the fabric’s H-block placement and aspect ratio; the *VTR toolchain* maps the circuit onto it. Given a netlist requiring n_{dsp} DSP and n_{bram} BRAM blocks, the policy emits

$$\pi = \left(\text{ar}, \{(x_i, y_i)\}_{i=1}^{n_{\text{dsp}}+n_{\text{bram}}} \right) : \quad (1)$$

a fabric aspect ratio $\text{ar} \in \mathcal{R} = \{0.1, 0.2, \dots, 2.0\}$, a discrete set of 20 candidate ratios spaced 0.1 apart, and a tile coordinate for every H-block. VTR then synthesizes, packs, places, routes, and auto-sizes the array to the smallest grid of the chosen aspect ratio that holds the design. The policy never places a CLB or fixes grid dimensions; it shapes the fabric and positions H-blocks; the cost is realized through VTR reports. We take that cost to be the three-term ADP (Section 1):

$$\text{ADP}(\pi) = \text{Area}(\pi) \times \text{Delay}(\pi) \times \text{Power}(\pi), \quad (2)$$

where each term is from the VTR flow (Section 4). Area is VPR’s *routing area*: since the number of H-blocks placed is fixed by functional requirements, logic block count offers limited sensitivity to placement decisions, while routing area varies directly with H-block positions through their influence on interconnect length and congestion, making it a more discriminative metric for placement quality. Delay is VPR’s *critical-path delay*: the longest timing path through the design’s logic and routing interconnects. Power is VPR’s *total power*: the dynamic and static power dissipated by the synthesized circuit. We report results as percentage ADP reduction relative to a traditional island-style baseline architecture. The Area-Delay-Power Product (ADP) acts as a comprehensive PPA metric, evaluating the design’s trade-offs by combining all three physical constraints into a single, joint optimization figure of merit.

3.2 RL Formulation

We pose placement as a finite-horizon sequential decision process and train with Maskable PPO [13] (MaskablePPO), which restricts the action space to legal tile coordinates at each step.

Episode structure. Step 0 selects the fabric aspect ratio from $\mathcal{R}$; each subsequent step positions one H-block (BRAMs then DSPs, both ordered by descending shared-net connectivity). Once all blocks are placed, the environment renders the layout into an architecture XML via Jinja2 and invokes VTR to obtain $\text{Area}(\pi)$, $\text{Delay}(\pi)$, $\text{Power}(\pi)$. Once these metrics are obtained, the final reward is computed by the environment and given to the agent intermediate steps receive zero reward.

Action space. The action space is Discrete($W_{\text{max}} \times H_{\text{max}}$), fixed across all benchmark circuits, where W_{max} and H_{max} are the core grid width and height taken over both the training pool and the unseen, zero-shot benchmark circuits (Section ??). Each benchmark’s own *core grid* is its width $\times$ height tile array excluding the surrounding IO ring, sized from its traditional-baseline resource requirements. At step 0, action index a selects $\mathcal{R}[a]$ directly (legal only for $a < |\mathcal{R}|$). At every later steps, the same index decodes to a tile coordinate $x = 1 + \lfloor a/H_{\text{max}} \rfloor$, $y = 1 + (a \bmod H_{\text{max}})$ within the IO ring, using the fixed H_{max} rather than the active benchmark circuit’s real height, so a given index always maps to the same (x, y) regardless of which benchmark circuit is active. Action masking then restricts the policy to indices whose decoded (x, y) falls inside the *active* benchmark circuit’s own core grid, fits the block’s height without crossing that grid’s boundary, and does not overlap an already-placed block; coordinates outside the active benchmark circuit’s core grid, even though addressable within the shared $W_{\text{max}} \times H_{\text{max}}$ space, are masked out every step.

Reward. The terminal reward is a sum of negative log-ratios against the traditional baseline for the active benchmark circuit:

$$r = - \left(\log \frac{\text{Area}(\pi)}{\text{Area}_0} + \log \frac{\text{Delay}(\pi)}{\text{Delay}_0} + \log \frac{\text{Power}(\pi)}{\text{Power}_0} \right), \quad (3)$$

where $\text{Area}_0, \text{Delay}_0, \text{Power}_0$ are the traditional-baseline metrics. An invalid choice (degenerate aspect ratio, out-of-bounds or overlapping placement, or a VTR failure) terminates early with penalty -10 .

3.3 Multi-Modal Feature Extractor

To generalize across benchmark circuits with varying netlist sizes, H-block counts, and grid dimensions, we encode each netlist as a graph, giving the policy the inductive bias for cross-topology transfer that graph learning provides elsewhere in EDA [21]. We reduce each netlist to a compact fixed-schema graph: one node per H-block instance plus one aggregate “fabric” node for the remaining standard-cell logic, with node features [is_dsp , is_bram , is_fabric , $\text{net_count_normalized}$]. Edges are weighted by shared-net count, normalized to $[0, 1]$, both between H-blocks and between each H-block and the fabric node.

The graph is padded to a fixed maximum node and edge count (Section ??) and encoded by two width-64 graph convolution layers, giving every H-block a learned embedding. Global mean-pooling yields a whole-netlist embedding, and the embedding of the node under consideration at the current step yields a per-step embedding; both are concatenated with a CNN encoding of the 4-channel occupancy grid and the benchmark circuit’s normalized fabric dimensions, then projected to a single 128-dimensional vector consumed by the PPO actor and critic heads. The encoder trains end-to-end under the PPO objective, with no separate pretraining stage.




Figure 1: System overview of the RL-based fabric layout search. The left panel encodes the input netlist and grid. The center panel rolls out the placement sequentially, positioning one H-block per step under action masking. The right panel uses Jinja2 to render the final placement into a VTR architecture XML for verification and reward computation.

3.4 Training Setup

Training samples uniformly from the benchmark pool, episode-by-episode. We use MaskablePPO with 24 parallel environments, 48 rollout steps per update, batch size 144, 15 epochs per update, and entropy coefficient 0.01. Each run trains for 13,000 episodes. We report three independent seeds (7, 42, 123) throughout Section 5 to characterize reliability rather than present a single run as ground truth.

4 Experimental Setup

4.1 Evaluation Oracle

Every candidate placement is scored by the true toolchain, not a proxy. A placement is rendered into a VTR architecture XML and VPR constraints file via Jinja2 templates, then run through the full VTR flow (synthesis, packing, placement, routing, and power estimation against a 45nm PTM file), and routing area, critical-path delay, and dynamic power are extracted from VPR reports.² Evaluated pairs are cached; the synthesis-to-routing flow dominates cost (seconds to low minutes per call).

4.2 Benchmark Circuits

Figure 2 lists all 17 benchmark circuits with grid size and H-block counts: 11 for training and 6 held-out for zero-shot evaluation only (Section 5.3), drawn from the VTR suite [18], Koios [1], and two small probes authored for this work, covering the small-fabric end the standard suites underrepresent. Four held-out benchmark circuits (softmax, reduction_layer, arm_core, mkDelayWorker32B) are larger circuits, probing fabric sizes well beyond the training




Figure 2: Per-seed ADP reduction for all 17 benchmark circuits (11 in-pool, 6 held-out zero-shot).

pool. The GA comparison (Section 5.2) uses softmax, held out of RL training.

5 Results

5.1 In-Pool Performance

Figure 2 (upper group) shows per-benchmark-circuit in-pool ADP reduction across three seeds. The pooled aggregate (summed ADP over all 11 benchmark circuits versus summed baseline) is 23.11%, 35.73%, and 23.97% for seeds 7, 42, and 123 respectively, averaging 27.60%. Nine of eleven benchmark circuits are consistently positive

²All evaluations use VTR 9.0.0 (git 22a09d39, GCC 13.3.0).

and stable; raygentop is the only persistent failure ($[-2.2, -1.7]\%$ across all seeds). or1200 and boundtop are positive but highly seed-sensitive, spanning 28–29 points each and driving most of the aggregate variance. In-pool performance is considerably more seed-sensitive than zero-shot, which stays within ~ 12 points across all six held-out benchmark circuits (Section 5.3): early convergence on most training benchmark circuits leaves fewer episodes for slow-to-converge ones like or1200, so which seed handles them well varies run to run.

5.2 GA Comparison

For a fair point of comparison, we implement a GA under the same oracle, ADP metric, and baseline on softmax: roulette-wheel selection, single-point crossover, per-gene mutation over H-block coordinates and aspect ratio, population 8, capped at 500 generations with patience 50. The GA and the policy search in fundamentally different ways, a GA generation is not an RL episode, so we compare the one quantity that is directly comparable between them: the best layout each one finds. The GA’s best reaches **14.62%** ADP reduction. Zero-shot, the policy already achieves **5.39–10.77%** across seeds (7, 42, 123) with no search at all; fine-tuned on softmax from seed 42, it reaches **14.82%**, matching the GA’s best.

The GA pays its search cost again from scratch for every new design. The trained policy pays its cost once and reuses it: zero-shot lands a positive result before any softmax-specific search begins, and the advantage compounds with the number of designs the same checkpoint is later applied to.

Figure 3 shows the physical saving on softmax. The traditional fabric provisions BRAM columns the netlist never uses; the policy omits them and selects an aspect ratio under which VTR auto-sizes the logic to a 22% smaller array, the dominant ADP term. The policy never moves a CLB; it removes idle columns and reshapes the fabric, and the tighter routing is VTR’s.

5.3 Zero-Shot Generalization

Figure 2 (lower group) reports ADP reduction on six held-out benchmark circuits across the same three seeds, with zero benchmark-specific training. All six fit within the checkpoint’s fixed dimensions. Only two (softmax, reduction_layer) were used to choose those dimensions; the other four were sized in by benchmark circuits elsewhere in the sizing superset, a harder test of generalization. Four of the six are larger than anything in the training pool (up to 2,304 tiles vs. 676), making this size extrapolation, not interpolation. Every benchmark circuit is positive under every seed (no sign flips), and every per-seed range is narrow (≤ 12 points), far from the or1200/boundtop spread. The overall average is **31.23%**.

6 Conclusion

We train one policy across 11 circuits and compare it to a GA under the same oracle, metric, and baseline (Section 5.2): the fine-tuned policy reaches a similar best ADP reduction, and zero-shot gives a positive result with no per-design search. The same policy applies zero-shot to six held-out benchmark circuits, four larger than anything in training (Section 5.3).

These results rest on three seeds and one deterministic rollout per zero-shot benchmark circuit, with in-pool fitting showing more seed

sensitivity than zero-shot transfer within that sample (Section 5.1). A few directions follow naturally: a systematic stochastic-rollout evaluation; extending the action space, currently placement and aspect ratio only, to co-evolve CLB LUT size and block sizing [3]; and running on the full GOLDS/NSGA-II benchmark circuit set for direct comparison.

References

- [1] Aman Arora, Andrew Boutros, Daniel Rauch, Aishwarya Rajen, Aatman Borda, Seyed A. Damghani, Samidh Mehta, Sangram Kate, Pragnesh Patel, Kenneth B. Kent, Vaughn Betz, and Lizy K. John. 2021. Koios: A Deep Learning Benchmark Suite for FPGA Architecture and CAD Research. In *31st International Conference on Field-Programmable Logic and Applications (FPL)*. IEEE. doi:10.1109/FPL53798.2021.00010
- [2] Vaughn Betz and Jonathan Rose. 1997. VPR: A New Packing, Placement and Routing Tool for FPGA Research. In *7th International Workshop on Field-Programmable Logic and Applications (FPL)*. Springer-Verlag, 213–222. doi:10.1007/3-540-63465-7_226
- [3] D. V. Bhargav, G. Pradyumna, and Madhav Rao. 2025. Meta-Heuristic Optimization of Custom Heterogeneous Blocks Defined eFPGA Design. In *38th International Conference on VLSI Design (VLSID)*. IEEE, 326–331. doi:10.1109/VLSID64188.2025.00069
- [4] Ruichen Chen, Shengyao Lu, Mohamed A. Elgammal, Peter Chun, Vaughn Betz, and Di Niu. 2023. VPR-Gym: A Platform for Exploring AI Techniques in FPGA Placement Optimization. In *33rd International Conference on Field-Programmable Logic and Applications (FPL)*. IEEE, 72–78. doi:10.1109/FPL60245.2023.00018
- [5] Ruoyu Cheng and Junchi Yan. 2021. On Joint Learning for Solving Placement and Routing in Chip Design. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 34. 16508–16519.
- [6] Luca Collini, Jitendra Bhandari, Chiara Muscarelli Tomajoli, Abdul Khader Thaklakkattu Moosa, Benjamin Tan, Xifan Tang, Pierre-Emmanuel Gaillardon, Ramesh Karri, and Christian Pilato. 2025. ARIANNA: An Automatic Design Flow for Fabric Customization and eFPGA Redaction. *ACM Transactions on Design Automation of Electronic Systems (TODAES)* (2025). volume/issue/pages unverified at press time. doi:10.1145/3737287
- [7] Seyed Alireza Damghani and Kenneth B. Kent. 2022. Yosys+Odin-II: The Odin-II Partial Mapper with Yosys Coarse-grained Netlists in VTR. In *Proceedings of the ACM/SIGDA International Symposium on Field-Programmable Gate Arrays (FPGA)*. 157. doi:10.1145/3490422.3502344
- [8] Voktho Das, Kimia Zamiri Azar, and Hadi M. Kamali. 2026. NuRedact: Non-Uniform eFPGA Architecture for Low-Overhead and Secure IP Redaction. In *Design, Automation & Test in Europe Conference (DATE)*. IEEE.
- [9] Mohamed A. Elgammal, Amin Mohaghegh, Soheil G. Shahrouz, Fatemehsadat Mahmoudi, Fahrican Kosar, Kimia Talaei, Joshua Fife, Daniel Khadivi, Kevin E. Murray, Andrew Boutros, Kenneth B. Kent, Jeffrey B. Goeters, and Vaughn Betz. 2025. VTR 9: Open-Source CAD for Fabric and Beyond FPGA Architecture Exploration. *ACM Transactions on Reconfigurable Technology and Systems (TRETS)* 18, 3 (2025), 39:1–39:53. doi:10.1145/3734798
- [10] Mohamed A. Elgammal, Kevin E. Murray, and Vaughn Betz. 2022. RLPlace: Using Reinforcement Learning and Smart Perturbations to Optimize FPGA Placement. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems (TCAD)* 41, 8 (2022), 2532–2545. doi:10.1109/TCAD.2021.3109863
- [11] Yunbo Hou, Haoran Ye, Shuwen Yang, Yingxue Zhang, Siyuan Xu, and Guojie Song. 2025. TransPlace: Transferable Circuit Global Placement via Graph Neural Network. *arXiv preprint arXiv:2501.05667* (2025).
- [12] Hao-Hsiang Hsiao, Yi-Chen Lu, Pruek Vanna-lampikul, and Sung Kyu Lim. 2024. FastTuner: Transferable Physical Design Parameter Optimization using Fast Reinforcement Learning. In *International Symposium on Physical Design (ISPD)*. ACM.
- [13] Shengyi Huang and Santiago Ontañón. 2022. A Closer Look at Invalid Action Masking in Policy Gradient Algorithms. In *Proceedings of the 35th International FLAIRS Conference*, Vol. 35. doi:10.32473/flairs.v35i.130584
- [14] Hadi M. Kamali, Kimia Z. Azar, Farimah Farahmandi, and Mark Tehranipoor. 2023. SheLL: Shrinking eFPGA Fabrics for Logic Locking. In *Design, Automation & Test in Europe Conference (DATE)*. IEEE, 1–6.
- [15] Yao Lai, Jinxin Liu, Zhentao Tang, Bin Wang, Jianye Hao, and Ping Luo. 2023. ChiFormer: Transferable Chip Placement via Offline Decision Transformer. In *Proceedings of the 40th International Conference on Machine Learning (ICML) (PMLR)*, Vol. 202. 18346–18364.
- [16] Yao Lai, Yao Mu, and Ping Luo. 2022. MaskPlace: Fast Chip Placement via Reinforced Visual Representation Learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 35. 24019–24030.
- [17] Azalia Mirhoseini, Anna Goldie, Mustafa Yazgan, Joe Wenjie Jiang, Ebrahim Songhori, Shen Wang, Young-Joon Lee, Eric Johnson, Omkar Pathak, Azade Nova, Jiwoo Pak, Andy Tong, Kavya Srinivasa, William Hang, Emre Tuncer,




Figure 3: What the policy recovers, physically (softmax, the GA comparison benchmark circuit). Both fabrics are drawn at the same tile scale. The traditional island-style baseline (left) is a 43×43 heterogeneous fabric that provisions BRAM columns (plum) the netlist never uses, with DSP blocks (terracotta) spread across regular DSP columns. The RL-optimized fabric (right) tailors the architecture to the netlist: it drops the unused BRAM columns entirely and repositions the DSPs, under an aspect ratio that lets VTR auto-size the same logic down to 36×40 . The dashed outline is the traditional footprint and the hatched strip is the reclaimed area: a 22% smaller fabric, the dominant term in the 15% ADP reduction. Gaps between tiles are routing channels.

- Quoc V. Le, James Laudon, Richard Ho, Roger Carpenter, and Jeff Dean. 2021. A graph placement methodology for fast chip design. *Nature* 594, 7862 (2021), 207–212. doi:10.1038/s41586-021-03544-w
- [18] Kevin E. Murray, Oleg Petelin, Sheng Zhong, Jia Min Wang, Mohamed Eldafrawy, Jean-Philippe Legault, Eugene Sha, Aaron G. Graham, Jean Wu, Matthew J. P. Walker, Hanqing Zeng, Panagiotis Patros, Jason Luu, Kenneth B. Kent, and Vaughn Betz. 2020. VTR 8: High-performance CAD and Customizable FPGA Architecture Modelling. *ACM Transactions on Reconfigurable Technology and Systems (TRETS)* 13, 2 (2020), 9:1–9:55. doi:10.1145/3388617
- [19] Pratyush Nandi, Anubhav Mishra, and Madhav Rao. 2024. GOLDS: Genetic Algorithm-based Optimization of Custom FPGA Architecture Layout Design for Secure Silicon. In *Proceedings of the Great Lakes Symposium on VLSI (GLSVLSI)*. ACM, 92–97. doi:10.1145/3649476.3658743
- [20] Chiara Muscari Tomajoli, Luca Collini, Jitendra Bhandari, Abdul Khader Thakkattu Moosa, Benjamin Tan, Xifan Tang, Pierre-Emmanuel Gaillardon, Ramesh Karri, and Christian Pilato. 2022. ALICE: An Automatic Design Flow for eFPGA Redaction. In *59th ACM/IEEE Design Automation Conference (DAC)*. 781–786. doi:10.1145/3489517.3530543
- [21] Nan Wu et al. 2023. GAMORA: Graph Learning based Symbolic Reasoning for Large-Scale Boolean Networks. In *IEEE/ACM International Conference on Computer-Aided Design (ICCAD)*.